

tugas-4

November 22, 2024

0.1 # Tugas 3 Data Mining - EDA + Klasifikasi

List Anggota Kelompok 9 * Cahaya Aulia Firdausyah (2006304) * Anderfa Jalu Kawani (2102671)
* Sabila Rosad (2106000)

0.2 Pendahuluan

Analisis Data Eksplorasi (EDA) dan klasifikasi memainkan peran penting dalam memahami pola historis dan membuat prediksi akurat. Pada kompetisi DMC tahun ini, task pertama berfokus pada analisis data sesi historis toko untuk memprediksi kemungkinan terjadinya pesanan (order). Data historis mencakup sekitar 50.000 sesi dengan atribut target “order” yang memiliki dua nilai: “y” untuk pesanan yang dilakukan dan “n” untuk tidak ada pesanan.

Melalui EDA, analisis awal terhadap data dilakukan untuk mengidentifikasi distribusi, pola, dan hubungan antar fitur yang berkontribusi pada kemungkinan pesanan. Informasi ini membantu membangun model prediktif berbasis machine learning untuk menentukan probabilitas [0,1] bagi setiap sesi baru. Evaluasi model dilakukan berdasarkan tingkat kesalahan terhadap hasil aktual dari sekitar 5.000 sesi yang disediakan. Dengan pendekatan ini, diharapkan solusi yang akurat dapat dihasilkan untuk mengoptimalkan prediksi pesanan di masa depan.

0.3 Import Library

Import library yang akan digunakan

```
[ ]: %matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab import drive
```

##Load Dataset

```
[ ]: # Mount Google Drive to Colab
drive.mount('/content/drive')

# Read the transact_train.txt file from Google Drive into a DataFrame
df = pd.read_csv("/content/drive/MyDrive/Data Mining/Tugas 3/transact_train.
↳txt", delimiter="|")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: # Display the first 5 rows of the DataFrame to preview the data
df.head()
```

```
[ ]: sessionNo  startHour  startWeekday  duration  cCount  cMinPrice  cMaxPrice  \
0           1           6           5      0.000        1      59.99      59.99
1           1           6           5     11.940        1      59.99      59.99
2           1           6           5     39.887        1      59.99      59.99
3           2           6           5      0.000        0         ?         ?
4           2           6           5     15.633        0         ?         ?

      cSumPrice  bCount  bMinPrice  ...      availability  customerNo  maxVal  \
0       59.99        1      59.99  ...                ?            1      600
1       59.99        1      59.99  ...  completely orderable            1      600
2       59.99        1      59.99  ...  completely orderable            1      600
3         ?         0         ?  ...  completely orderable            ?         ?
4         ?         0         ?  ...  completely orderable            ?         ?

      customerScore  accountLifetime  payments  age  address  lastOrder  order
0                70                21         1  43        1         49        y
1                70                21         1  43        1         49        y
2                70                21         1  43        1         49        y
3                 ?                 ?         ?  ?        ?         ?         y
4                 ?                 ?         ?  ?        ?        ?         y

[5 rows x 24 columns]
```

```
[ ]: # Show a summary of the DataFrame, including column data types, non-null
      counts, and memory usage
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 429013 entries, 0 to 429012
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sessionNo             429013 non-null  int64
1   startHour              429013 non-null  int64
2   startWeekday           429013 non-null  int64
3   duration                429013 non-null  float64
4   cCount                  429013 non-null  int64
5   cMinPrice               429013 non-null  object
6   cMaxPrice               429013 non-null  object
7   cSumPrice               429013 non-null  object
8   bCount                  429013 non-null  int64
9   bMinPrice               429013 non-null  object
```

```

10 bMaxPrice      429013 non-null object
11 bSumPrice      429013 non-null object
12 bStep          429013 non-null object
13 onlineStatus   429013 non-null object
14 availability    429013 non-null object
15 customerNo     429013 non-null object
16 maxVal         429013 non-null object
17 customerScore  429013 non-null object
18 accountLifetime 429013 non-null object
19 payments       429013 non-null object
20 age           429013 non-null object
21 address        429013 non-null object
22 lastOrder      429013 non-null object
23 order          429013 non-null object
dtypes: float64(1), int64(5), object(18)
memory usage: 78.6+ MB

```

```

[ ]: # Generate summary statistics (mean, median, standard deviation, etc.) for
      ↪ numerical columns
df.describe()

```

```

[ ]:
count    sessionNo    startHour    startWeekday    duration \
mean    25274.631293    14.617061    5.924839    1573.901640
std      14441.366146    4.485914    0.790930    2427.123356
min       1.000000    0.000000    5.000000    0.000000
25%     12731.000000    11.000000    5.000000    225.070000
50%     25470.000000    15.000000    6.000000    738.199000
75%     37542.000000    18.000000    7.000000    1880.265000
max      50000.000000    23.000000    7.000000    21580.092000

count    cCount    bCount
mean      24.140317    4.135168
std       30.398164    4.451778
min        0.000000    0.000000
25%        5.000000    1.000000
50%       13.000000    3.000000
75%       31.000000    5.000000
max       200.000000   108.000000

```

```

[ ]: # Identify columns containing "?"
columns_with_question_mark = [col for col in df.columns if df[col].astype(str).
      ↪ str.contains('\?').any()]

print("Columns containing '?':", columns_with_question_mark)

```

Columns containing '?': ['cMinPrice', 'cMaxPrice', 'cSumPrice', 'bMinPrice',

```
'bMaxPrice', 'bSumPrice', 'bStep', 'onlineStatus', 'availability', 'customerNo',
'maxVal', 'customerScore', 'accountLifetime', 'payments', 'age', 'address',
'lastOrder']
```

```
[ ]: # Convert numeric columns to float, handling '?' values
numeric_cols = ['sessionNo', 'startHour', 'startWeekday', 'duration', 'cCount',
↳ 'bCount', 'maxVal', 'customerScore', 'accountLifetime', 'payments', 'age',
↳ 'lastOrder']
for col in numeric_cols:
    # Replace '?' with NaN before converting to float
    df[col] = df[col].replace('?', np.nan).astype(float)

# Convert categorical columns to string
categorical_cols = ['cMinPrice', 'cMaxPrice', 'cSumPrice', 'bMinPrice',
↳ 'bMaxPrice', 'bSumPrice', 'bStep', 'onlineStatus', 'availability', 'order',
↳ 'customerNo', 'address']
df[categorical_cols] = df[categorical_cols].astype(str)

# Verify the data types
print(df.dtypes)
```

```
sessionNo      float64
startHour      float64
startWeekday   float64
duration       float64
cCount         float64
cMinPrice      object
cMaxPrice      object
cSumPrice      object
bCount         float64
bMinPrice      object
bMaxPrice      object
bSumPrice      object
bStep          object
onlineStatus   object
availability   object
customerNo     object
maxVal         float64
customerScore  float64
accountLifetime float64
payments       float64
age            float64
address        object
lastOrder      float64
order          object
dtype: object
```

```
[ ]: # Define a list of columns that are numeric (i.e., continuous or quantitative
      ↪ values)
numeric_cols = ['sessionNo', 'startHour', 'startWeekday', 'duration', 'cCount',
      ↪ 'bCount', 'maxVal', 'customerScore', 'accountLifetime', 'payments', 'age',
      ↪ 'lastOrder']

# Define a list of columns that are categorical (i.e., discrete or qualitative
      ↪ values)
categorical_cols = ['cMinPrice', 'cMaxPrice', 'cSumPrice', 'bMinPrice',
      ↪ 'bMaxPrice', 'bSumPrice', 'bStep', 'onlineStatus', 'availability', 'order',
      ↪ 'customerNo', 'address']

[ ]: from sklearn.impute import SimpleImputer

# Impute missing values in numeric columns using mean
numeric_imputer = SimpleImputer(strategy='mean')
df[numeric_cols] = numeric_imputer.fit_transform(df[numeric_cols])

[ ]: # Impute missing values in categorical columns using mode
from collections import Counter

for col in categorical_cols:
    mode_value = df[col].mode().iloc[0]
    df[col] = df[col].fillna(mode_value)

[ ]: # Check for remaining '?' values
print(df.isin(['?']).sum())
```

sessionNo	0
startHour	0
startWeekday	0
duration	0
cCount	0
cMinPrice	2765
cMaxPrice	2765
cSumPrice	2765
bCount	0
bMinPrice	5130
bMaxPrice	5130
bSumPrice	5130
bStep	191333
onlineStatus	160379
availability	165255
customerNo	151098
maxVal	0
customerScore	0
accountLifetime	0

```

payments          0
age               0
address          151098
lastOrder         0
order             0
dtype: int64

```

```

[ ]: # Check for NaN values
      print(df.isna().sum())

```

```

sessionNo        0
startHour        0
startWeekday     0
duration         0
cCount          0
cMinPrice        0
cMaxPrice        0
cSumPrice        0
bCount          0
bMinPrice        0
bMaxPrice        0
bSumPrice        0
bStep           0
onlineStatus     0
availability     0
customerNo       0
maxVal          0
customerScore    0
accountLifetime  0
payments        0
age             0
address         0
lastOrder       0
order           0
dtype: int64

```

```

[ ]: # Display the first 5 rows of the DataFrame to preview the data
      df.head()

```

```

[ ]:   sessionNo  startHour  startWeekday  duration  cCount  cMinPrice  cMaxPrice  \
0         1.0         6.0         5.0     0.000      1.0      59.99      59.99
1         1.0         6.0         5.0    11.940      1.0      59.99      59.99
2         1.0         6.0         5.0    39.887      1.0      59.99      59.99
3         2.0         6.0         5.0     0.000      0.0         ?         ?
4         2.0         6.0         5.0    15.633      0.0         ?         ?

      cSumPrice  bCount  bMinPrice  ...      availability  customerNo  \

```

0	59.99	1.0	59.99	...	?	1
1	59.99	1.0	59.99	...	completely orderable	1
2	59.99	1.0	59.99	...	completely orderable	1
3	?	0.0	?	...	completely orderable	?
4	?	0.0	?	...	completely orderable	?

	maxVal	customerScore	accountLifetime	payments	age	address \
0	600.00000	70.000000	21.000000	1.000000	43.000000	1
1	600.00000	70.000000	21.000000	1.000000	43.000000	1
2	600.00000	70.000000	21.000000	1.000000	43.000000	1
3	2486.35827	485.298449	135.557403	15.218016	44.919861	?
4	2486.35827	485.298449	135.557403	15.218016	44.919861	?

	lastOrder	order
0	49.000000	y
1	49.000000	y
2	49.000000	y
3	79.883975	y
4	79.883975	y

[5 rows x 24 columns]

```
[ ]: # Summary statistics for numeric columns
print(df[numeric_cols].describe())

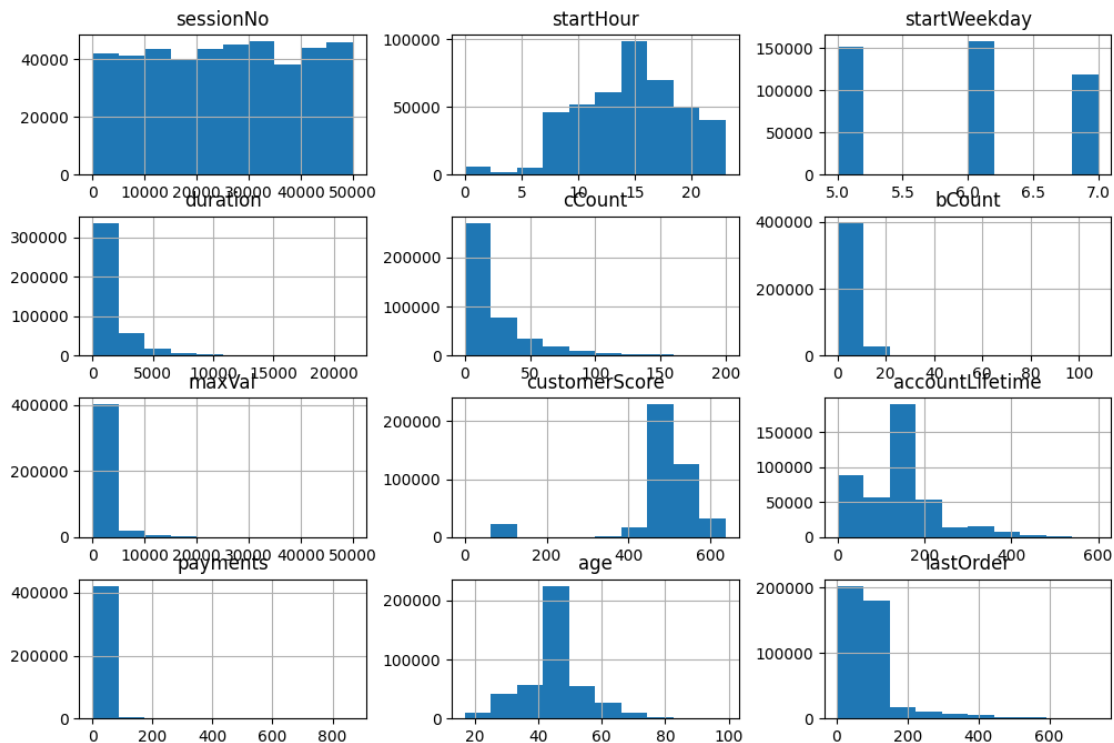
# Visualize the distributions of numeric columns
import matplotlib.pyplot as plt
df[numeric_cols].hist(figsize=(12, 8))
plt.show()
```

	sessionNo	startHour	startWeekday	duration \
count	429013.000000	429013.000000	429013.000000	429013.000000
mean	25274.631293	14.617061	5.924839	1573.901640
std	14441.366146	4.485914	0.790930	2427.123356
min	1.000000	0.000000	5.000000	0.000000
25%	12731.000000	11.000000	5.000000	225.070000
50%	25470.000000	15.000000	6.000000	738.199000
75%	37542.000000	18.000000	7.000000	1880.265000
max	50000.000000	23.000000	7.000000	21580.092000

	cCount	bCount	maxVal	customerScore \
count	429013.000000	429013.000000	429013.000000	429013.000000
mean	24.140317	4.135168	2486.358270	485.298449
std	30.398164	4.451778	2433.856317	104.956462
min	0.000000	0.000000	0.000000	0.000000
25%	5.000000	1.000000	900.000000	485.298449
50%	13.000000	3.000000	2486.358270	485.298449

75%	31.000000	5.000000	2500.000000	535.000000
max	200.000000	108.000000	50000.000000	638.000000

	accountLifetime	payments	age	lastOrder
count	429013.000000	429013.000000	429013.000000	429013.000000
mean	135.557403	15.218016	44.919861	79.883975
std	87.774074	28.083922	9.601616	91.111767
min	0.000000	0.000000	17.000000	3.000000
25%	75.000000	6.000000	41.000000	23.000000
50%	135.557403	15.218016	44.919861	79.883975
75%	156.000000	15.218016	48.000000	79.883975
max	600.000000	868.000000	99.000000	738.000000



```
[ ]: # Explore the categorical columns
for col in categorical_cols:
    print(f"Column: {col}")

    # Print the top 10 most frequent values
    value_counts = df[col].value_counts()
    print("Most frequent values:")
    print(value_counts.head(10))

    # Check for any unexpected or invalid values
```



```

unusual_values = value_counts[value_counts < 10].index
if len(unusual_values) > 0:
    print("Unusual/Infrequent values:")
    print(unusual_values)
print("----")

```

Column: cMinPrice

Most frequent values:

cMinPrice

9.99	55667
3.99	35395
19.99	23352
12.99	19007
14.99	18227
29.99	15337
4.99	15227
6.99	13211
7.99	12986
24.99	11050

Name: count, dtype: int64

Unusual/Infrequent values:

```

Index(['65.6', '34.96', '119.9', '46.99', '19.96', '569.99', '98.0', '110.0',
      '48.74', '18.5',
      ...,
      '37.49', '31.95', '888.0', '23.5', '51.99', '159.9', '1190.0', '519.0',
      '30.95', '54.9'],
      dtype='object', name='cMinPrice', length=263)

```

Column: cMaxPrice

Most frequent values:

cMaxPrice

29.99	33194
19.99	27636
49.99	26032
39.99	24772
24.99	17868
59.99	16131
99.99	12498
79.99	11288
59.95	10973
34.99	10499

Name: count, dtype: int64

Unusual/Infrequent values:

```

Index(['28.9', '28.5', '247.12', '165.99', '42.5', '185.0', '57.85', '859.99',
      '449.0', '675.0',
      ...,
      '34.96', '18.16', '11.0', '1190.0', '35.96', '23.5', '95.99', '4.75',

```

```

        '15.3', '201.99'],
        dtype='object', name='cMaxPrice', length=245)
---
Column: cSumPrice
Most frequent values:
cSumPrice
49.99      3942
?          2765
6.99       2754
39.98      2461
59.98      1692
89.97      1686
19.98      1602
59.97      1577
79.98      1553
49.98      1368
Name: count, dtype: int64
Unusual/Infrequent values:
Index(['605.65', '749.57', '270.88', '314.79', '526.63', '1651.81', '1674.84',
      '172.88', '919.88', '367.81',
      ...
      '836.53', '900.49', '167.24', '198.15', '228.13', '258.12', '2275.91',
      '553.86', '650.61', '5253.28'],
      dtype='object', name='cSumPrice', length=64510)
---
Column: bMinPrice
Most frequent values:
bMinPrice
9.99       56187
3.99       32670
19.99      27620
14.99      21855
12.99      20079
29.99      19388
24.99      13461
39.99      11275
49.99      11187
6.99       10548
Name: count, dtype: int64
Unusual/Infrequent values:
Index(['1039.99', '1.45', '134.99', '82.49', '36.0', '366.0', '67.49', '32.0',
      '57.99', '1.25',
      ...
      '51.97', '13.5', '36.76', '37.96', '7.8', '549.95', '43.0', '62.96',
      '9.79', '159.9'],
      dtype='object', name='bMinPrice', length=239)
---
Column: bMaxPrice

```

Most frequent values:

bMaxPrice

29.99	40203
19.99	36399
39.99	25841
49.99	25754
24.99	21253
59.99	14910
9.99	14629
14.99	13085
59.95	10684
34.99	10018

Name: count, dtype: int64

Unusual/Infrequent values:

```
Index(['6.59', '71.95', '909.9', '2799.99', '28.5', '26.95', '2199.99',
      '387.03', '67.49', '6999.99',
      ...
      '4.65', '154.9', '598.0', '62.0', '69.74', '349.95', '299.9', '1619.0',
      '1549.99', '7.8'],
      dtype='object', name='bMaxPrice', length=216)
```

Column: bSumPrice

Most frequent values:

bSumPrice

29.99	8733
19.99	8285
49.99	8082
39.99	6129
9.99	5879
24.99	5413
?	5130
14.99	4005
39.98	3886
59.98	3489

Name: count, dtype: int64

Unusual/Infrequent values:

```
Index(['228.52', '324.38', '2264.0', '212.94', '421.76', '151.97', '1291.59',
      '452.81', '397.86', '161.61',
      ...
      '49.36', '102.83', '122.82', '241.82', '160.81', '210.76', '235.75',
      '273.69', '772.96', '211.86'],
      dtype='object', name='bSumPrice', length=15207)
```

Column: bStep

Most frequent values:

bStep

?	191333
1	90058

```

2      60682
4      41142
3      30062
5      15736
Name: count, dtype: int64
---
Column: onlineStatus
Most frequent values:
onlineStatus
y      265625
?      160379
n       3009
Name: count, dtype: int64
---
Column: availability
Most frequent values:
availability
completely orderable      253692
?      165255
mainly orderable          5756
completely not orderable  1491
mixed                    1284
completely not determinable 1017
mainly not orderable      320
mainly not determinable   198
Name: count, dtype: int64
---
Column: order
Most frequent values:
order
y      290030
n      138983
Name: count, dtype: int64
---
Column: customerNo
Most frequent values:
customerNo
?      151098
47      1248
5464     345
4953     196
16740    155
7394     152
5488     149
4034     136
15769    123
4118     118
Name: count, dtype: int64

```

```

Unusual/Infrequent values:
Index(['15891', '12761', '22717', '14906', '14882', '12456', '22729', '12736',
      '14524', '20945',
      ...
      '11690', '11665', '11494', '11626', '11570', '11544', '11540', '11537',
      '11512', '10461'],
      dtype='object', name='customerNo', length=13625)

```

```

---
Column: address
Most frequent values:
address
2    203570
?    151098
1     74058
3       287
Name: count, dtype: int64
---

```

```

[ ]: # Cross-validate columns
print(df.loc[df['cCount'] > df['bCount']])

```

	sessionNo	startHour	startWeekday	duration	cCount	cMinPrice	\
8	3.0	6.0	5.0	181.477	9.0	29.99	
9	3.0	6.0	5.0	297.018	11.0	9.99	
10	3.0	6.0	5.0	310.967	11.0	9.99	
11	3.0	6.0	5.0	324.278	11.0	9.99	
12	3.0	6.0	5.0	341.613	11.0	9.99	
...	
429006	49998.0	18.0	7.0	2961.909	6.0	59.99	
429007	49998.0	18.0	7.0	4700.383	50.0	9.99	
429008	49998.0	18.0	7.0	5988.882	77.0	9.99	
429009	49999.0	18.0	7.0	675.114	6.0	59.0	
429010	49999.0	18.0	7.0	715.341	7.0	59.0	

	cMaxPrice	cSumPrice	bCount	bMinPrice	...	availability	\
8	29.99	89.97	1.0	29.99	...	?	
9	29.99	109.95	2.0	9.99	...	?	
10	29.99	109.95	2.0	9.99	...	completely orderable	
11	29.99	109.95	2.0	9.99	...	completely orderable	
12	29.99	109.95	2.0	9.99	...	completely orderable	
...	
429006	99.99	419.94	1.0	59.99	...	?	
429007	119.99	2974.6	2.0	59.99	...	?	
429008	149.99	5253.28	3.0	49.95	...	?	
429009	199.99	509.96	1.0	89.99	...	?	
429010	649.99	1159.95	1.0	89.99	...	completely orderable	

	customerNo	maxVal	customerScore	accountLifetime	payments	\
--	------------	--------	---------------	-----------------	----------	---

8	3	1800.00000	475.000000	302.000000	12.000000
9	3	1800.00000	475.000000	302.000000	12.000000
10	3	1800.00000	475.000000	302.000000	12.000000
11	3	1800.00000	475.000000	302.000000	12.000000
12	3	1800.00000	475.000000	302.000000	12.000000
...
429006	?	2486.35827	485.298449	135.557403	15.218016
429007	?	2486.35827	485.298449	135.557403	15.218016
429008	?	2486.35827	485.298449	135.557403	15.218016
429009	25038	2486.35827	485.298449	135.557403	0.000000
429010	25038	2486.35827	485.298449	135.557403	0.000000

	age	address	lastOrder	order
8	45.000000	1	11.000000	y
9	45.000000	1	11.000000	y
10	45.000000	1	11.000000	y
11	45.000000	1	11.000000	y
12	45.000000	1	11.000000	y
...
429006	44.919861	?	79.883975	n
429007	44.919861	?	79.883975	n
429008	44.919861	?	79.883975	n
429009	24.000000	1	4.000000	n
429010	24.000000	1	4.000000	n

[373604 rows x 24 columns]

```
[ ]: # Identify outliers using z-score
from scipy.stats import zscore

z = np.abs(zscore(df[numeric_cols]))

# Create a boolean mask for outlier rows, considering any outlier across columns
outlier_mask = (z > 3).any(axis=1)

# Filter the DataFrame using the outlier mask
outliers = df[outlier_mask]

print(outliers)
```

	sessionNo	startHour	startWeekday	duration	cCount	cMinPrice	\
0	1.0	6.0	5.0	0.000	1.0	59.99	
1	1.0	6.0	5.0	11.940	1.0	59.99	
2	1.0	6.0	5.0	39.887	1.0	59.99	
70	12.0	6.0	5.0	9.220	2.0	9.99	
71	12.0	6.0	5.0	91.283	4.0	5.99	
...	
428999	49996.0	18.0	7.0	7015.682	197.0	6.99	

429000	49996.0	18.0	7.0	7074.729	197.0	6.99
429001	49996.0	18.0	7.0	7089.360	197.0	6.99
429002	49996.0	18.0	7.0	7170.905	197.0	6.99
429003	49996.0	18.0	7.0	7271.812	197.0	6.99


	cMaxPrice	cSumPrice	bCount	bMinPrice	...	availability	\
0	59.99	59.99	1.0	59.99	...	?	
1	59.99	59.99	1.0	59.99	...	completely orderable	
2	59.99	59.99	1.0	59.99	...	completely orderable	
70	9.99	19.98	1.0	9.99	...	?	
71	9.99	31.96	2.0	5.99	...	?	
...	
428999	59.99	4315.03	13.0	9.99	...	completely orderable	
429000	59.99	4315.03	13.0	9.99	...	completely orderable	
429001	59.99	4315.03	13.0	9.99	...	completely orderable	
429002	59.99	4315.03	13.0	9.99	...	completely orderable	
429003	59.99	4315.03	13.0	9.99	...	completely orderable	


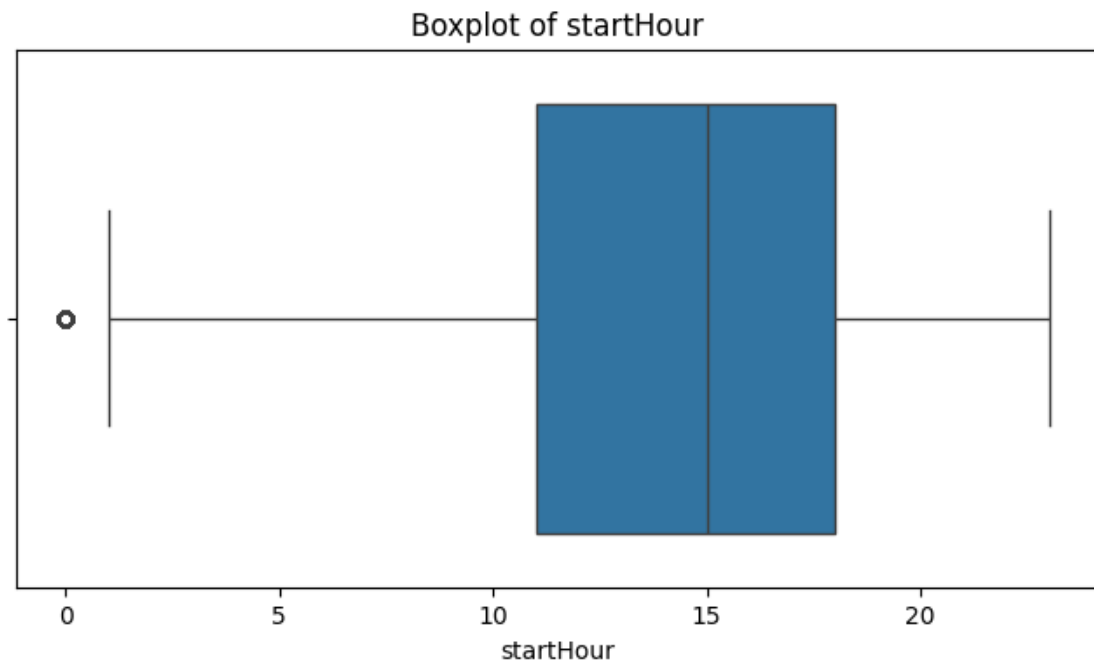
	customerNo	maxVal	customerScore	accountLifetime	payments	\
0	1	600.00000	70.000000	21.000000	1.000000	
1	1	600.00000	70.000000	21.000000	1.000000	
2	1	600.00000	70.000000	21.000000	1.000000	
70	8	2000.00000	546.000000	364.000000	11.000000	
71	8	2000.00000	546.000000	364.000000	11.000000	
...	
428999	?	2486.35827	485.298449	135.557403	15.218016	
429000	?	2486.35827	485.298449	135.557403	15.218016	
429001	?	2486.35827	485.298449	135.557403	15.218016	
429002	?	2486.35827	485.298449	135.557403	15.218016	
429003	?	2486.35827	485.298449	135.557403	15.218016	

	age	address	lastOrder	order
0	43.000000	1	49.000000	y
1	43.000000	1	49.000000	y
2	43.000000	1	49.000000	y
70	86.000000	2	37.000000	y
71	86.000000	2	37.000000	y
...
428999	44.919861	?	79.883975	y
429000	44.919861	?	79.883975	y
429001	44.919861	?	79.883975	y
429002	44.919861	?	79.883975	y
429003	44.919861	?	79.883975	y


[69869 rows x 24 columns]

```
[ ]: # Boxplots for each numeric column to spot outliers
for col in numeric_cols:
    plt.figure(figsize=(8, 4))
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    plt.show()
```







Boxplot of duration




Boxplot of cCount




Boxplot of bCount




Boxplot of maxVal


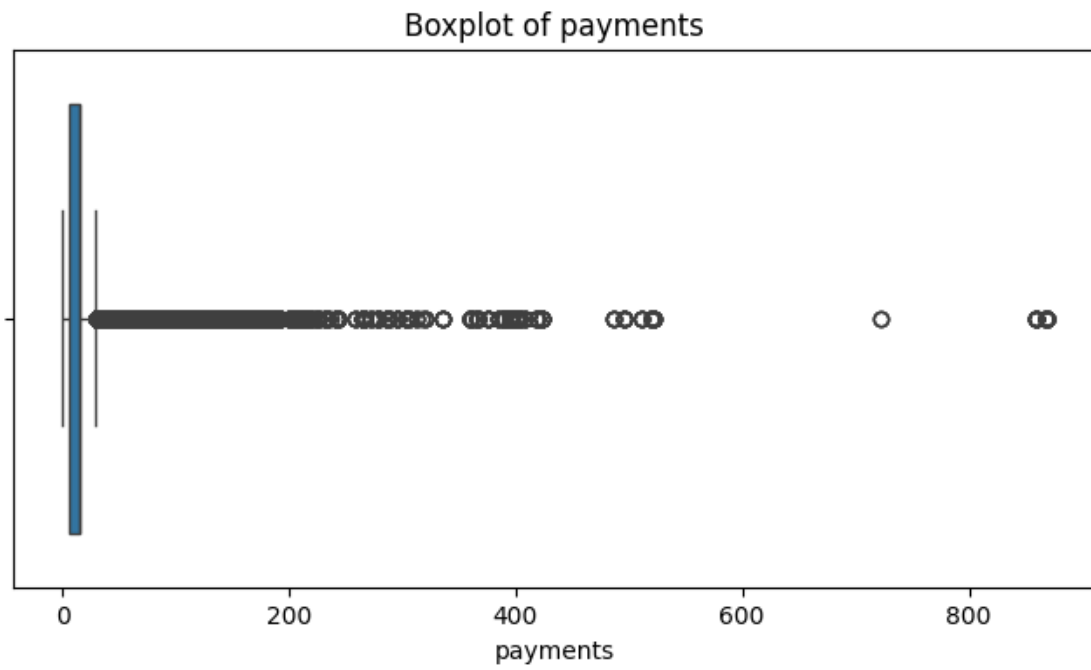



Boxplot of customerScore



Boxplot of accountLifetime







```
[ ]: #Convert all to lowercase
df[categorical_cols] = df[categorical_cols].apply(lambda x: x.str.lower())
```

```
[ ]: # Check logical consistency between min and max price columns
inconsistent_prices = df[df['cMinPrice'] > df['cMaxPrice']]
print("Rows with inconsistent price values:\n", inconsistent_prices)
```

Rows with inconsistent price values:

	sessionNo	startHour	startWeekday	duration	cCount	cMinPrice	\
9	3.0	6.0	5.0	297.018	11.0	9.99	
10	3.0	6.0	5.0	310.967	11.0	9.99	
11	3.0	6.0	5.0	324.278	11.0	9.99	
12	3.0	6.0	5.0	341.613	11.0	9.99	
24	6.0	6.0	5.0	0.000	2.0	99.99	
...	
429002	49996.0	18.0	7.0	7170.905	197.0	6.99	
429003	49996.0	18.0	7.0	7271.812	197.0	6.99	
429007	49998.0	18.0	7.0	4700.383	50.0	9.99	
429008	49998.0	18.0	7.0	5988.882	77.0	9.99	
429009	49999.0	18.0	7.0	675.114	6.0	59.0	

	cMaxPrice	cSumPrice	bCount	bMinPrice	...	availability	\
9	29.99	109.95	2.0	9.99	...	?	
10	29.99	109.95	2.0	9.99	...	completely orderable	
11	29.99	109.95	2.0	9.99	...	completely orderable	

12	29.99	109.95	2.0	9.99	...	completely	orderable
24	129.99	229.98	2.0	99.99	...		?
...		
429002	59.99	4315.03	13.0	9.99	...	completely	orderable
429003	59.99	4315.03	13.0	9.99	...	completely	orderable
429007	119.99	2974.6	2.0	59.99	...		?
429008	149.99	5253.28	3.0	49.95	...		?
429009	199.99	509.96	1.0	89.99	...		?

	customerNo	maxVal	customerScore	accountLifetime	payments	\
9	3	1800.00000	475.000000	302.000000	12.000000	
10	3	1800.00000	475.000000	302.000000	12.000000	
11	3	1800.00000	475.000000	302.000000	12.000000	
12	3	1800.00000	475.000000	302.000000	12.000000	
24	?	2486.35827	485.298449	135.557403	15.218016	
...	
429002	?	2486.35827	485.298449	135.557403	15.218016	
429003	?	2486.35827	485.298449	135.557403	15.218016	
429007	?	2486.35827	485.298449	135.557403	15.218016	
429008	?	2486.35827	485.298449	135.557403	15.218016	
429009	25038	2486.35827	485.298449	135.557403	0.000000	

	age	address	lastOrder	order
9	45.000000	1	11.000000	y
10	45.000000	1	11.000000	y
11	45.000000	1	11.000000	y
12	45.000000	1	11.000000	y
24	44.919861	?	79.883975	n
...
429002	44.919861	?	79.883975	y
429003	44.919861	?	79.883975	y
429007	44.919861	?	79.883975	n
429008	44.919861	?	79.883975	n
429009	24.000000	1	4.000000	n

[150046 rows x 24 columns]

```
[ ]: # Correct inconsistent price values
df.loc[df['cMinPrice'] > df['cMaxPrice'], ['cMinPrice', 'cMaxPrice']] = df.
    ↪loc[df['cMinPrice'] > df['cMaxPrice'], ['cMaxPrice', 'cMinPrice']].values

[ ]: # Display summary statistics for all columns in the DataFrame, including both
    ↪numeric and categorical columns
print(df.describe(include='all'))
```

	sessionNo	startHour	startWeekday	duration	\
count	429013.000000	429013.000000	429013.000000	429013.000000	
unique	NaN	NaN	NaN	NaN	

top	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN
mean	25274.631293	14.617061	5.924839	1573.901640
std	14441.366146	4.485914	0.790930	2427.123356
min	1.000000	0.000000	5.000000	0.000000
25%	12731.000000	11.000000	5.000000	225.070000
50%	25470.000000	15.000000	6.000000	738.199000
75%	37542.000000	18.000000	7.000000	1880.265000
max	50000.000000	23.000000	7.000000	21580.092000

	cCount	cMinPrice	cMaxPrice	cSumPrice	bCount	bMinPrice	\
count	429013.000000	429013	429013	429013	429013.000000	429013	
unique	NaN	839	825	72990	NaN	748	
top	NaN	19.99	9.99	49.99	NaN	9.99	
freq	NaN	36762	57229	3942	NaN	56187	
mean	24.140317	NaN	NaN	NaN	4.135168	NaN	
std	30.398164	NaN	NaN	NaN	4.451778	NaN	
min	0.000000	NaN	NaN	NaN	0.000000	NaN	
25%	5.000000	NaN	NaN	NaN	1.000000	NaN	
50%	13.000000	NaN	NaN	NaN	3.000000	NaN	
75%	31.000000	NaN	NaN	NaN	5.000000	NaN	
max	200.000000	NaN	NaN	NaN	108.000000	NaN	

	...	availability	customerNo	maxVal	customerScore	\
count	...	429013	429013	429013.000000	429013.000000	
unique	...	8	25038	NaN	NaN	
top	...	completely	orderable	?	NaN	NaN
freq	...	253692	151098	NaN	NaN	
mean	...	NaN	NaN	2486.358270	485.298449	
std	...	NaN	NaN	2433.856317	104.956462	
min	...	NaN	NaN	0.000000	0.000000	
25%	...	NaN	NaN	900.000000	485.298449	
50%	...	NaN	NaN	2486.358270	485.298449	
75%	...	NaN	NaN	2500.000000	535.000000	
max	...	NaN	NaN	50000.000000	638.000000	

	accountLifetime	payments	age	address	lastOrder	\
count	429013.000000	429013.000000	429013.000000	429013	429013.000000	
unique	NaN	NaN	NaN	4	NaN	
top	NaN	NaN	NaN	2	NaN	
freq	NaN	NaN	NaN	203570	NaN	
mean	135.557403	15.218016	44.919861	NaN	79.883975	
std	87.774074	28.083922	9.601616	NaN	91.111767	
min	0.000000	0.000000	17.000000	NaN	3.000000	
25%	75.000000	6.000000	41.000000	NaN	23.000000	
50%	135.557403	15.218016	44.919861	NaN	79.883975	
75%	156.000000	15.218016	48.000000	NaN	79.883975	
max	600.000000	868.000000	99.000000	NaN	738.000000	

	order
count	429013
unique	2
top	y
freq	290030
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

[11 rows x 24 columns]

Model Klasifikasi

```
[ ]: import pandas as pd #Make sure pandas is imported
from sklearn.model_selection import train_test_split #Import train_test_split
    ↳to split the data
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
    ↳confusion_matrix

# Assuming 'df' is your DataFrame containing all the data

X = df.drop(columns=['order'])
y = df['order']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↳random_state=42)

# Initialize LabelEncoder
encoder = LabelEncoder()

# Iterate through columns of X_train and encode object (string) types
for col in X_train.select_dtypes(include=['object']).columns:
    # Fit on the combined unique values from both training and testing data
    all_values = pd.concat([X_train[col], X_test[col]]).unique()
    encoder.fit(all_values)

    X_train[col] = encoder.transform(X_train[col])
    X_test[col] = encoder.transform(X_test[col]) # Apply the same encoding to
    ↳X_test
```

```

# Create and train the RandomForestClassifier
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)

# Make predictions on the test data #This line is added to get predictions
↳from the model
y_pred_rf = rf_model.predict(X_test)

# Calculate accuracy
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print("Accuracy (Random Forest):", accuracy_rf)

# show classification reports and confusion matrix
print("Classification Report (Random Forest):\n", classification_report(y_test,
↳y_pred_rf))
print("Confusion Matrix (Random Forest):\n", confusion_matrix(y_test,
↳y_pred_rf))

```

Accuracy (Random Forest): 0.9122524853443352

Classification Report (Random Forest):

	precision	recall	f1-score	support
n	0.89	0.83	0.86	27700
y	0.92	0.95	0.94	58103
accuracy			0.91	85803
macro avg	0.91	0.89	0.90	85803
weighted avg	0.91	0.91	0.91	85803

Confusion Matrix (Random Forest):

```

[[22983  4717]
 [ 2812 55291]]

```

0.4 Analisis Model Klasifikasi

1. Akurasi:

- Akurasi model adalah 0.9122, yang berarti model memprediksi kelas dengan benar sekitar 91.22% dari waktu.
- Ini merupakan nilai akurasi yang cukup tinggi, menunjukkan bahwa model secara umum berkinerja baik.

2. Classification Report:

- Precision:
 - Precision untuk kelas “n” adalah 0.89, yang berarti dari semua data yang diprediksi sebagai “n”, sekitar 89% benar-benar “n”.
 - Precision untuk kelas “y” adalah 0.92, yang berarti dari semua data yang diprediksi sebagai “y”, sekitar 92% benar-benar “y”.

- Recall:
 - Recall untuk kelas “n” adalah 0.83, yang berarti model berhasil mengidentifikasi sekitar 83% dari semua data “n” yang sebenarnya.
 - Recall untuk kelas “y” adalah 0.95, yang berarti model berhasil mengidentifikasi sekitar 95% dari semua data “y” yang sebenarnya.
- F1-score:
 - F1-score adalah rata-rata harmonik antara precision dan recall.
 - F1-score untuk kelas “n” adalah 0.86, dan untuk kelas “y” adalah 0.94.
 - Nilai F1-score yang tinggi menunjukkan keseimbangan yang baik antara precision dan recall.
- Support:
 - Support menunjukkan jumlah data aktual untuk setiap kelas.
 - Ada 27700 data untuk kelas “n” dan 58103 data untuk kelas “y”.

3. Confusion Matrix:


- Confusion matrix memberikan gambaran lebih detail tentang kinerja model:
 - True Positive (TP): 55291 (model memprediksi “y” dengan benar)
 - True Negative (TN): 22983 (model memprediksi “n” dengan benar)
 - False Positive (FP): 4717 (model memprediksi “y” secara salah)
 - False Negative (FN): 2812 (model memprediksi “n” secara salah)

0.5 Visualisasi Confusion Matrix

- Predicted No, Actual No (True Negative - TN): 22,983. Model memprediksi “No” dengan benar untuk 22,983 sampel.
- Predicted Yes, Actual No (False Positive - FP): 4,717. Model salah memprediksi “Yes” padahal seharusnya “No”.
- Predicted No, Actual Yes (False Negative - FN): 2,812. Model salah memprediksi “No” padahal seharusnya “Yes”.
- Predicted Yes, Actual Yes (True Positive - TP): 55,291. Model memprediksi “Yes” dengan benar untuk 55,291 sampel.

```
[ ]: cm = confusion_matrix(y_test, y_pred_rf)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
             xticklabels=['Predicted No', 'Predicted Yes'],
             yticklabels=['Actual No', 'Actual Yes'])
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```



0.6 Visualisasi ROC Curve

- Dengan AUC sebesar 0.97, model memiliki kemampuan diskriminasi yang sangat baik untuk membedakan antara kelas positif dan negatif.
- Pada berbagai threshold, model mampu mempertahankan keseimbangan yang baik antara TPR (benar memprediksi positif) dan FPR (meminimalkan prediksi positif palsu).

```
[ ]: from sklearn.metrics import roc_curve, auc

# Assuming you have probabilities for the positive class (e.g., from
# predict_proba)
y_probs = rf_model.predict_proba(X_test)[: , 1]

# Convert 'y_test' to numerical format where 'y' is 1 and 'n' is 0
y_test_numeric = [1 if value == 'y' else 0 for value in y_test]


# Use the numerical y_test in roc_curve
fpr, tpr, thresholds = roc_curve(y_test_numeric, y_probs)
```

```

roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
        roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()

```



0.7 Visualisasi Feature Importance


- sessionNo, duration, dan cSumPrice memiliki nilai feature importance tertinggi. Hal ini menunjukkan bahwa fitur-fitur tersebut memiliki pengaruh paling besar dalam menentukan output model.

- availability, startWeekday, dan onlineStatus memiliki kontribusi yang sangat kecil terhadap prediksi model. Fitur ini mungkin tidak terlalu relevan atau penting bagi model.

```
[ ]: importances = rf_model.feature_importances_
feature_names = X_train.columns

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(10, 6))
plt.title("Feature Importance")
plt.bar(range(X_train.shape[1]), importances[indices], align="center")
plt.xticks(range(X_train.shape[1]), feature_names[indices], rotation=90)
plt.tight_layout()
plt.show()
```



0.8 Kesimpulan

1. **Akurasi Model:** Model ini memiliki akurasi **91.22%**, yang menunjukkan bahwa model memprediksi kelas dengan benar sekitar 91.22% dari waktu.
2. **Classification Report:**
 - **Precision:** Untuk kelas “n” (Tidak Order), precision mencapai **0.89**, dan untuk kelas “y” (Order), precision mencapai **0.92**.
 - **Recall:** Untuk kelas “n”, recall adalah **0.83**, sedangkan untuk kelas “y”, recall mencapai **0.95**.

- **F1-score:** Nilai F1-score untuk kelas “n” adalah **0.86**, sementara untuk kelas “y” adalah **0.94**.
- **Support:** Jumlah data untuk kelas “n” adalah **27,700**, sementara kelas “y” memiliki **58,103** data.

3. Confusion Matrix:

- Model menghasilkan **True Positive (TP)** sebanyak **55,291** untuk kelas “y”.
 - Model menghasilkan **True Negative (TN)** sebanyak **22,983** untuk kelas “n”.
 - Terdapat **4,717 False Positive (FP)**, di mana model salah memprediksi kelas “y” padahal seharusnya “n”.
 - Terdapat **2,812 False Negative (FN)**, di mana model salah memprediksi kelas “n” padahal seharusnya “y”.
4. **Visualisasi ROC Curve:** AUC (Area Under the Curve) model adalah **0.97**, menunjukkan bahwa model memiliki kemampuan diskriminasi yang sangat baik dalam membedakan antara kelas positif dan negatif.
5. **Visualisasi Feature Importance:** Fitur yang memiliki kontribusi tertinggi dalam model adalah **sessionNo**, **duration**, dan **cSumPrice**, sementara fitur dengan kontribusi lebih kecil adalah **availability**, **startWeekday**, dan **onlineStatus**.

Secara keseluruhan, model ini menunjukkan performa yang sangat baik, dengan akurasi tinggi, precision dan recall yang seimbang, serta kemampuan yang kuat dalam membedakan antara kelas positif dan negatif.